

b02902080 郭傳駿

b02902100 康耀文

# parallel & distributed computing

## final report

### Mancos

- MICRO ARCHITECTURE AND NETWORK CO-SIMULATOR
- TRACE PACKET FLOW ACTIVITY
- APPLICATION: KEEPING TRACK OF
  - NETWORK FLOW
  - TRAFFIC FLOW
  - CASH FLOW
- INPUT: LOG FILE WITH ATTRIBUTES

```
FGETS:96938, 2, 2, 3, 2, 1760215299, 3850983036, 4224, 35784, 5001, 11251, 1514, 96938
id:1549, name:net_recv, nid:2, nidsrc:3, niddst:2, seq:1760215299, ack:3850983036, flag:4224, portsrc:35784,
portdst:5001, insns:11251, payload:1514, lt:96938

FGETS:96939, 2, 2, 3, 2, 284213507, 3850983036, 4224, 35784, 5001, 0, 1514, 96939
id:1550, name:net_recv, nid:2, nidsrc:3, niddst:2, seq:284213507, ack:3850983036, flag:4224, portsrc:35784,
portdst:5001, insns:0, payload:1514, lt:96939

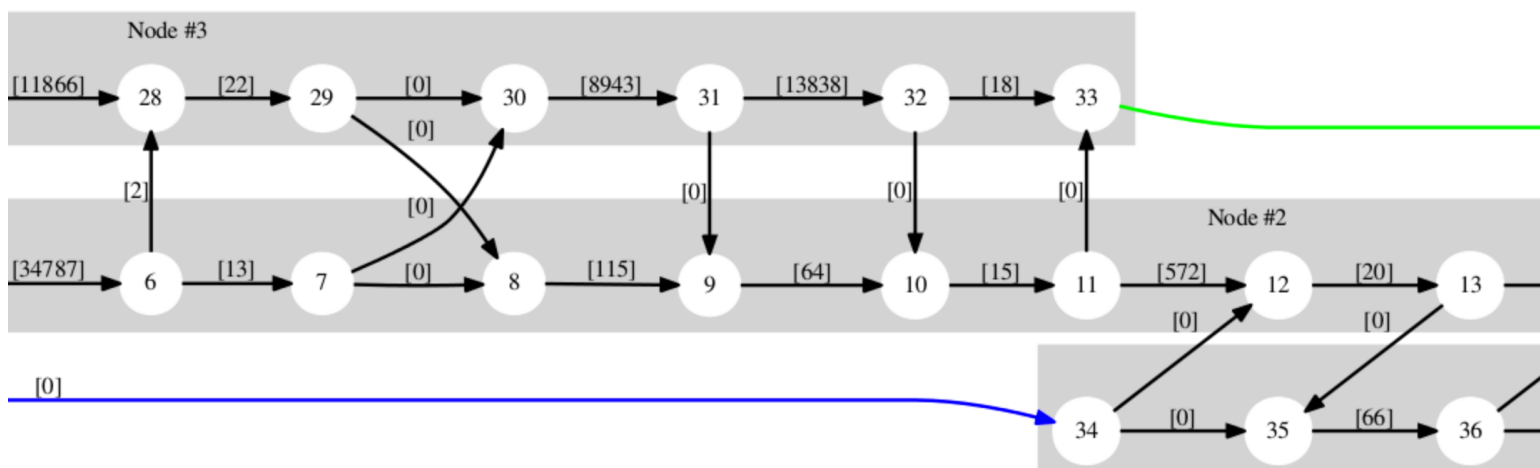
FGETS:96940, 2, 2, 3, 2, 3103113475, 3850983036, 4224, 35784, 5001, 0, 1514, 96940
id:1551, name:net_recv, nid:2, nidsrc:3, niddst:2, seq:3103113475, ack:3850983036, flag:4224, portsrc:35784,
portdst:5001, insns:0, payload:1514, lt:96940

FGETS:96941, 2, 2, 3, 2, 1627111683, 3850983036, 4224, 35784, 5001, 0, 1514, 96941
id:1552, name:net_recv, nid:2, nidsrc:3, niddst:2, seq:1627111683, ack:3850983036, flag:4224, portsrc:35784,
portdst:5001, insns:0, payload:1514, lt:96941

FGETS:96942, 2, 2, 3, 2, 134332931, 3850983036, 4224, 35784, 5001, 0, 1514, 96942
id:1553, name:net_recv, nid:2, nidsrc:3, niddst:2, seq:134332931, ack:3850983036, flag:4224, portsrc:35784,
portdst:5001, insns:0, payload:1514, lt:96942

FGETS:96943, 2, 2, 3, 2, 2953232899, 3850983036, 4224, 35784, 5001, 0, 1514, 96943
id:1554, name:net_recv, nid:2, nidsrc:3, niddst:2, seq:2953232899, ack:3850983036, flag:4224, portsrc:35784,
portdst:5001, insns:0, payload:1514, lt:96943
```

- OUTPUT: FLOW GRAPH

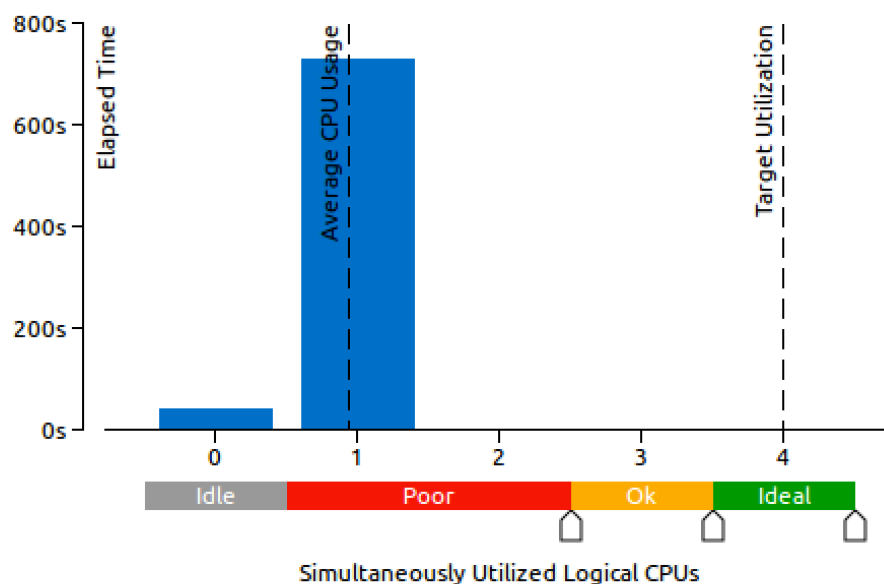


## Why do we parallelise it?

The application must be scalable in reasonable time in order to deal with real world network activity between thousands of end-devices(nodes).

## Vtune profile of the original version

Function Stack ▼	CPU Time: Total			
	Effective Time by Utilization			Spin Time
	Idle	Poor	Ok	
▼ Total	100.0%			0.0%
▼ start	100.0%			0.0%
▼ _libc_start_main	100.0%			0.0%
▼ main	100.0%			0.0%
▼ weaver_post_process	100.0%			0.0%
▼ std::vector<SimA...	0.0%			0.0%
▶ _ZNSt16allocator<char>	0.0%			0.0%
▼ longest_path	99.8%			0.0%
▶ std::_Destroy<Si...	0.0%			0.0%
▼ std::_uninitializ	57.5%			0.0%
▼ std::uninitializ	57.5%			0.0%
▼ std::uninitia	57.5%			0.0%
▼ std::copy<	57.5%			0.0%
▶ std::_mit	0.0%			0.0%
▼ std::_cop	57.5%			0.0%
▶ std::_nil	0.0%			0.0%
▼ std::_co	57.4%			0.0%
Selected 1 row(s):				100.0%



## main bottleneck:

- function find\_id\_of\_remote\_dest
- find activity pairs(send/receive) over all activities in all nodes
- requires  $O(N*N)$  Time  $N$ = total num of activity

```
uint32_t findIdOfRemoteDest( vector<SimActivity> saque, SimActivity sa ){
    uint32_t min_lt = UINT_MAX;
    SimActivity sa_cand;
    uint32_t id_ret = sa.id;

    //printf("Activity ID %u : ConnectToRemote, lt(%u)=%u, type=%s\n", sa.id,

    for( int i=0; i<saque.size(); i++ ){
        sa_cand = saque.at(i);
        //printf("sa.nid=%u, sa_cand : nid=%u id=%u lt(%u)=%u\n", sa.nid, sa_cand.nid, sa_cand.id, sa_cand.lt);
        if( (sa.id != sa_cand.id) && (sa_cand.nid == sa.nid_dst) && (sa.seq == sa_cand.seq) ){
            printf("%u connect to %u, s%u d%u ; s%u d%u, flag:%u %u ack: %u %u\n", sa.id, sa_cand.id, sa.nid, sa_cand.nid, sa.seq, sa_cand.seq, sa.flag, sa_cand.flag, sa.ack, sa_cand.ack);
            id_ret = sa_cand.id;
            break;
        }
    }

    return id_ret;
}
```

## code optimization:

### 1. KEEP TRACK OF EXISTING NODE AND FIRST ID OF EACH NODE

- int first\_id[10] = {1};
- int node\_id[10] = {0};
- if(flag == 0){  
first\_id[n] = global\_id;  
flag = 1;  
}

### 2. ACCORDING TO VTUNE ANALYSIS, THE COPYING OF "SAQUE" TAKES A LOT OF TIME

- we changed the function parameter from call by value to call by reference

### 3. IN FUNCTION FIND\_ID\_OF\_REMOTE\_DEST

- search for pairs only in the section of the destination node
- this also requires activities of each node to be in time order

```
uint32_t findIdOfRemoteDest( vector<SimActivity> &saque, SimActivity sa ){
    uint32_t min_lt = UINT_MAX;
    SimActivity sa_cand;
    uint32_t id_ret = sa.id;
    int ssize = saque.size();
    //printf("Activity ID %u : ConnectToRemote, lt(%u)=%u, type=%s\n", sa.id, sa.nid, sa.lt, type);

    for( int i=first_id[sa.nid_dst]-1; i<ssize; i++ ){
        sa_cand = saque.at(i);
        //printf("sa.nid=%u, sa_cand : nid=%u id=%u lt(%u)=%u\n", sa.nid, sa_cand.nid, sa_cand.id, sa_cand.lt);
        if( (sa.id != sa_cand.id) && (sa_cand.nid == sa.nid_dst) && (sa.seq == sa_cand.seq) && (sa.lt < sa_cand.lt) ){
            //printf("%u connect to %u, s%u d%u ; s%u d%u, flag:%u %u ack: %u %u\n", sa.id, sa_cand.id, sa.nid, sa_cand.nid, sa.seq, sa_cand.seq, sa.flag, sa_cand.flag, sa.ack, sa_cand.ack);
            id_ret = sa_cand.id;
            break;
        }
    }

    return id_ret;
}
```

#### 4. IN FUNCTION FIND\_ID\_OF\_LOCAL\_SRC

- actually, there is no need to find id of local source
- if the id belongs to the first activity of a node, the local source should be itself
- for other activities, the local source would be the previous activity

```
uint32_t findIdOfLocalSrc( vector<SimActivity> &saque, SimActivity sa ){  
    int flag = 0;  
    for(int n=0; n<MAX_NUM_NODES; n++)  
        if(sa.id == first_id[n])  
            flag = 1;  
    if(flag == 0)  
        return sa.id - 1;  
    else  
        return sa.id;  
}
```

#### 5. FIRST / LAST ACTIVITY OF NODES ARE ALREADY RECORDED

```
SimActivity findTheFirstActivityOfNode( vector<SimActivity> &saque, const uint32_t nid ){  
    uint32_t min_id = UINT_MAX;  
    SimActivity sa_ret = sa_begin;  
  
    if(node_id[nid] == 1){  
        //printf("%d\n",nid );  
        SimActivity sa = saque.at(first_id[nid] - 1);  
        sa_ret = sa;  
    }  
  
    return sa_ret;  
}
```

```
SimActivity findTheLastActivityOfNode( const vector<SimActivity> &saque, const uint32_t nid ){  
    uint32_t max_id = 0;  
    SimActivity sa_ret = sa_end;  
    int next = -1;  
  
    if(node_id[nid] == 1){  
        for( int i=nid+1; i < MAX_NUM_NODES; i++ ){  
            if(node_id[i] == 1){  
                next = i;  
                break;  
            }  
        }  
    }  
  
    if(next != -1){  
        //printf("%d\n",nid );  
        SimActivity sa = saque.at(first_id[next] - 2);  
        sa_ret = sa;  
    }  
  
    else if(node_id[nid] == 1 && next == -1){  
        //printf("%d\n",nid );  
        SimActivity sa = saque.at(saque.size() - 1);  
        sa_ret = sa;  
    }  
  
    return sa_ret;  
    //printf("Find last activity at node %u :\n", nid);  
}
```

## 6. VECTOR INDEXING BY OPERATOR[] IS FASTER THAN OPERATOR .AT()

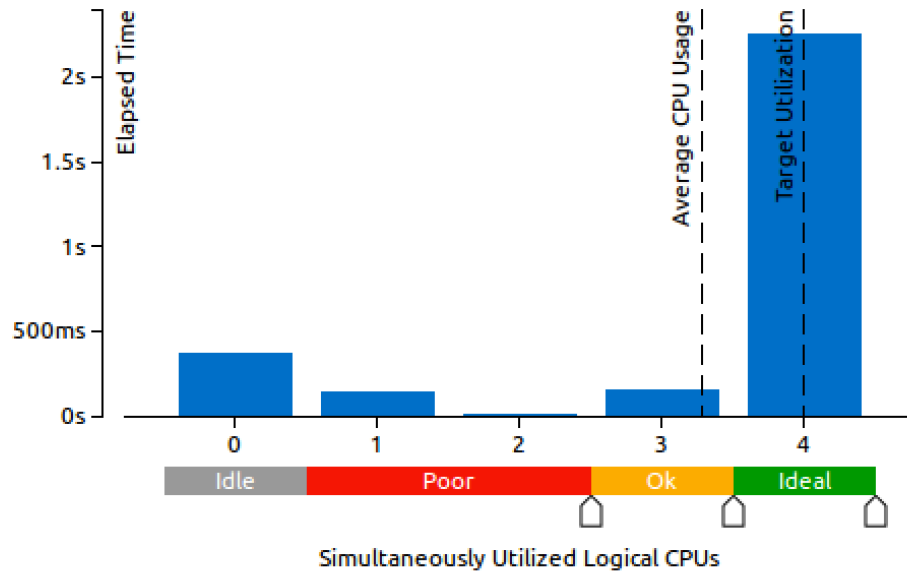
### parallelize with OpenMP

```
#pragma omp parallel
{
    #pragma omp for schedule(dynamic)
    for( i=0 ; i<saque.size() ; i++ ){
        SimActivity &sa = saque.at(i);
        if(sa.type == NETWORK_SEND){
            sa.id_rdst = findIdOfRemoteDest(saque, sa);
            //printf("S: Add edge from %u to %u : payload = %u\n", sa.id, sa.id_rdst, sa.payload );
            if( sa.id_rdst != sa.id ){
                //g_mtx.lock();
                g.addEdge( sa.id, sa.id_rdst, sa.nettime_us );//nettime unit : us
                //g_mtx.unlock();
            }
        }

        //printf("Before find local source\n");
        sa.id_ldst = findIdOfLocalSrc(saque, sa);
        //printf("Add edge from %u to %u : payload = %u\n", sa.id, sa.id_ldst, sa.payload );
        if( sa.id_ldst != sa.id ){
            //g_mtx.lock();
            g.addEdge( sa.id_ldst, sa.id, sa.cputime_us ); //cputime unit : us
            //g_mtx.unlock();
        }
    }
}
```

### analysis of our results

Grouping: Function / Call Stack		CPU Time			Module	
Function / Call Stack	Effective Time by Utilization	Spin Tim.	Ov.. Tim.			
	Idle Poor Ok Ideal Over					
▼ findIdOfRemoteDest	6.714s		0s	0s	mancos-gather	findIdOfRemoteDest
▼ longest_path_omp_fn.0	6.714s		0s	0s	mancos-gather	longest_path_omp_fn.0
func@0x8290 ← start_thread ← _clone	5.002s		0s	0s	libgomp.so.1	func@0x8290
longest_path ← weaver_post_process ← main ← __libc_start_main	1.712s		0s	0s	mancos-gather	longest_path(std::vector<SimActivity, std::allocator<SimActivity>>::operator[])
▶ std::vector<SimActivity, std::allocator<SimActivity>>::operator[]	2.742s		0s	0s	mancos-gather	std::vector<SimActivity, std::allocator<SimActivity>>::operator[]
▶ __sscanf	0.072s		0s	0s	libc.so.6	__sscanf
▶ longest_path	0.058s		0s	0s	mancos-gather	longest_path(std::vector<SimActivity, std::allocator<SimActivity>>::operator[])
▶ std::ostream::M_insert<long>	0.022s		0s	0s	libstdc++.so.6	std::ostream& std::ostream::M_insert<long>()
▶ func@0x9ae0	0.016s		0s	0s	libgomp.so.1	func@0x9ae0
▶ Graph::topologicalSortUtil	0.012s		0s	0s	mancos-gather	Graph::topologicalSortUtil
▶ std::List_base<AdjListNode, std::allocator<AdjListNode>>::push_back	0.012s		0s	0s	mancos-gather	std::List_base<AdjListNode, std::allocator<AdjListNode>>::push_back
▶ func@0x4f50	0.008s		0s	0s	libgomp.so.1	func@0x4f50
▶ __gnu_cxx::new_allocator<std::List_node<AdjListNode>>::construct	0.008s		0s	0s	mancos-gather	__gnu_cxx::new_allocator<std::List_node<AdjListNode>>::construct
▶ std::stack<int, std::deque<int, std::allocator<int>>>::pop	0.008s		0s	0s	mancos-gather	std::stack<int, std::deque<int, std::allocator<int>>>::pop
▶ findIdOfLocalSrc	0.008s		0s	0s	mancos-gather	findIdOfLocalSrc



## comparison

- CPU 2.4 GHZ INTEL CORE I5 TURBO UP TO 2.9GHZ
- 4 LOGICAL CPUS

### ORIGINAL

```
real    13m2.415s
user    12m7.263s
sys     0m8.526s
```

### OPTIMIZED & PARALLELIZED (4 THREADS)

```
real    0m10.799s
user    0m37.780s
sys     0m0.226s
```

### • SPEEDUP

- total speedup about 72x
- speedup due to multithreading 3.7/10 about 3.7
- trying on different number of threads

	1 thread	2 thread	3 thread	4 thread
real time	26.419	13.931	11.159	10.799
user time	26.233	27.285	32.569	37.78
theoretical speedup	~1	1.96	2.91	3.49

